Introduction to Data Science

UNIVERSITY OF
SAN FRANCISCO

James D. Wilson

MSAN 601 - Linear Regression Analysis

Outline

- Course Overview
- What is Data Science?
  - Where is Data Science?
  - A brief history

A Little About Me

- Ph.D. Statistics and Operations Research (UNC Chapel Hill, ’15)
  - Research focused on statistical analysis of networks
  - Explore, model, and analyze network data (e.g., social networks)
- M.S. Mathematical Sciences (Clemson University, ’10)
- B.S. Mathematics and Chemistry (Campbell University ’08)

A Little About Me

Classes I teach:

- BSDS 100 - Intro to Data Science with R
- MATH 106 - Business Statistics
- MATH 370 - Probability with Applications
- MATH 373 - Statistical Learning
- MSAN 601 - Linear Regression Analysis
- MSAN 630 - Advanced Computational Statistics
- MSAN 700 - Social Network Analysis

A Little About Me

- Born and raised in NC (near Raleigh)
- Live in Rockridge, Berkeley.
- A huge college basketball fan! (Go Heels!)
- Have loved college football since 2008 (Go Tigers!)
- Enjoy tasting beers (bourbon-barrel stouts are my favorite).

Course Description and Syllabus

All lecture notes, the syllabus, assignments, and course description are available at this course website:

https://github.com/jdwilson4/Regression-Analysis

What is Data Science?

- **Wikipedia**: “the extraction of knowledge from data."
- A precise definition is a bit unclear and has faced much controversy... (we’ll see more on this in a moment)
- Practitioners tend to agree on the *components* of data science:
  - gathering and cleaning data
  - database management
  - exploratory analysis
  - predictive modeling
  - data summary and visualization

Where is Data Science?

- Twitter feed, December 2014

The Evolution of Data Science

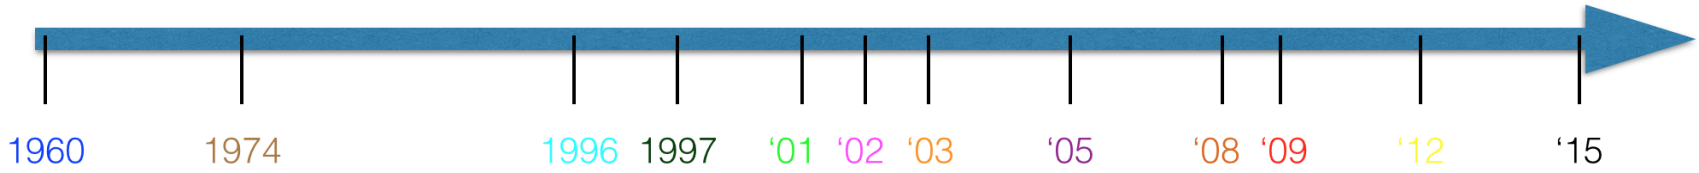

- 1960: Peter Naur (CS Ph.D.) published *Datalogy: the science of data and its place in education.*
- 1974: Peter Naur published *Concise Survey of Computer Methods*.
  - defines data science as “the science of dealing with data, once they have been established."
  - continues to say that “... the relation of the data to what they represent is delegated to other fields and sciences."

The Evolution of Data Science

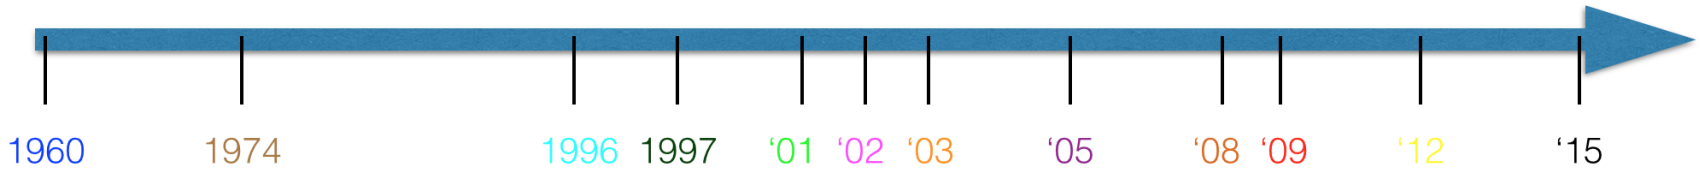

- 1996: International Federation of Classification Societies meet in Tokyo and for the first time include “data science" in the conference title: “Data science, classification, and related methods."

- 1997: C.F. Jeff Wu gave the inaugural lecture “Statistics = Data Science?" for appointment to the H. C. Carver Professorship at the University of Michigan.

The Evolution of Data Science

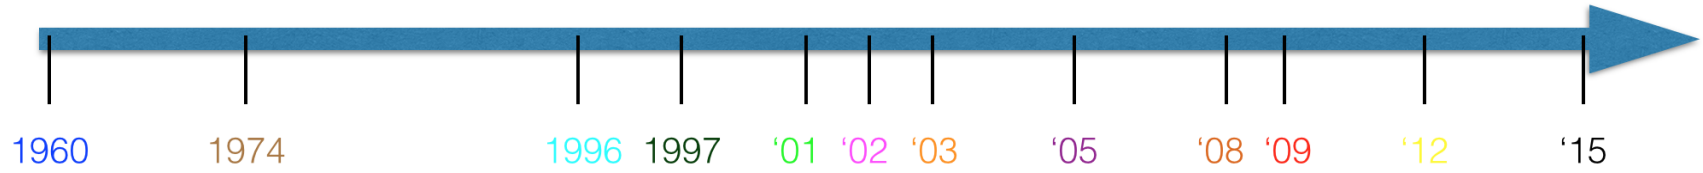

- 2001: William Cleveland (Bell Labs) published *Data Science: An Action Plan for Expanding the Technical Areas of the Field of Statistics*.
  - Sets forth 6 areas for a university department involving statistics.
- 2002: *Data Science Journal* is launched
  - Focus on data systems, publications on internet, and applications
- 2003: *Journal of Data Science* is launched
  - Focus on application of statistical and quantitative methods

The Evolution of Data Science

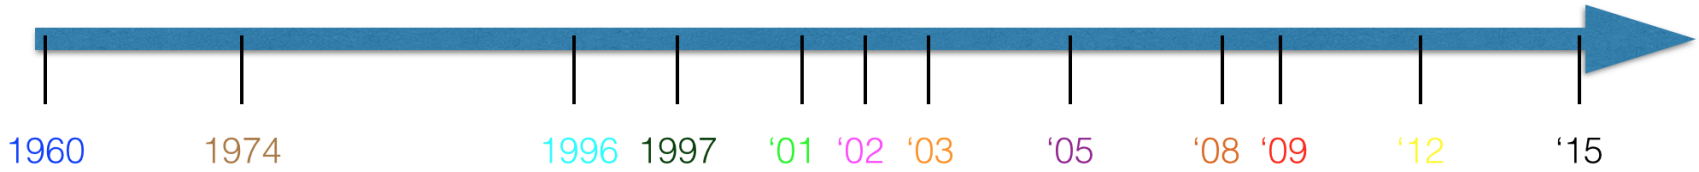

- 2005: National Science board redefines data scientists:
  - “The information and computer scientists, data and software programmers, disciplinary experts, ... who are crucial to successful management of a digital data collection whose primary activity is to conduct creative inquiry and analysis"
- 2008: DJ Patil (LinkedIn) and Jeff Hammerbacher (Facebook) coined the term "data scientist" to define their jobs

The Evolution of Data Science

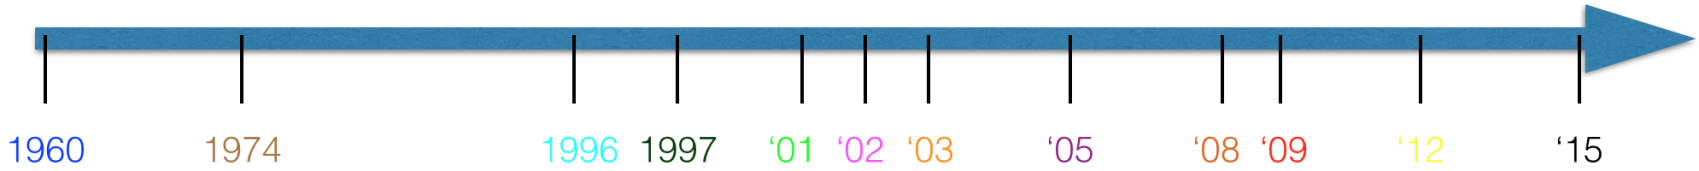

- January, 2009: Hal Varian (chief economist at Google) writes that "... the sexy job in the next 10 years will be statisticians."
- October, 2012: Harvard Business Review publishes "Data Scientist: The Sexiest Job of the 21st Century."
- February 5th, 2015: DJ Patil appointed as the first Chief Data Scientist in the White House.

Applications

Marketing analytics, sports analytics, biotechnology, social experiments, e-commerce, government analysis, ...

Why Data Science?

- Size, complexity, and amount of data
  - Predicted $\approx$ 40 trillion gigabytes of data in 2020; up from 130 billion in 2005!
  - Big data requires innovative techniques for analysis
- *McKinsey*: "The U.S. faces a shortage of 140K - 190K people with analytical expertise and 1.5 million managers and analysts with the skills to understand and make decisions based on the analysis of big data." (May, 2011)
- *Harvard Business Review*: "Data Scientist: The Sexiest Job of the 21st Century." (October, 2012)

Data Scientists: The unicorn industries want?

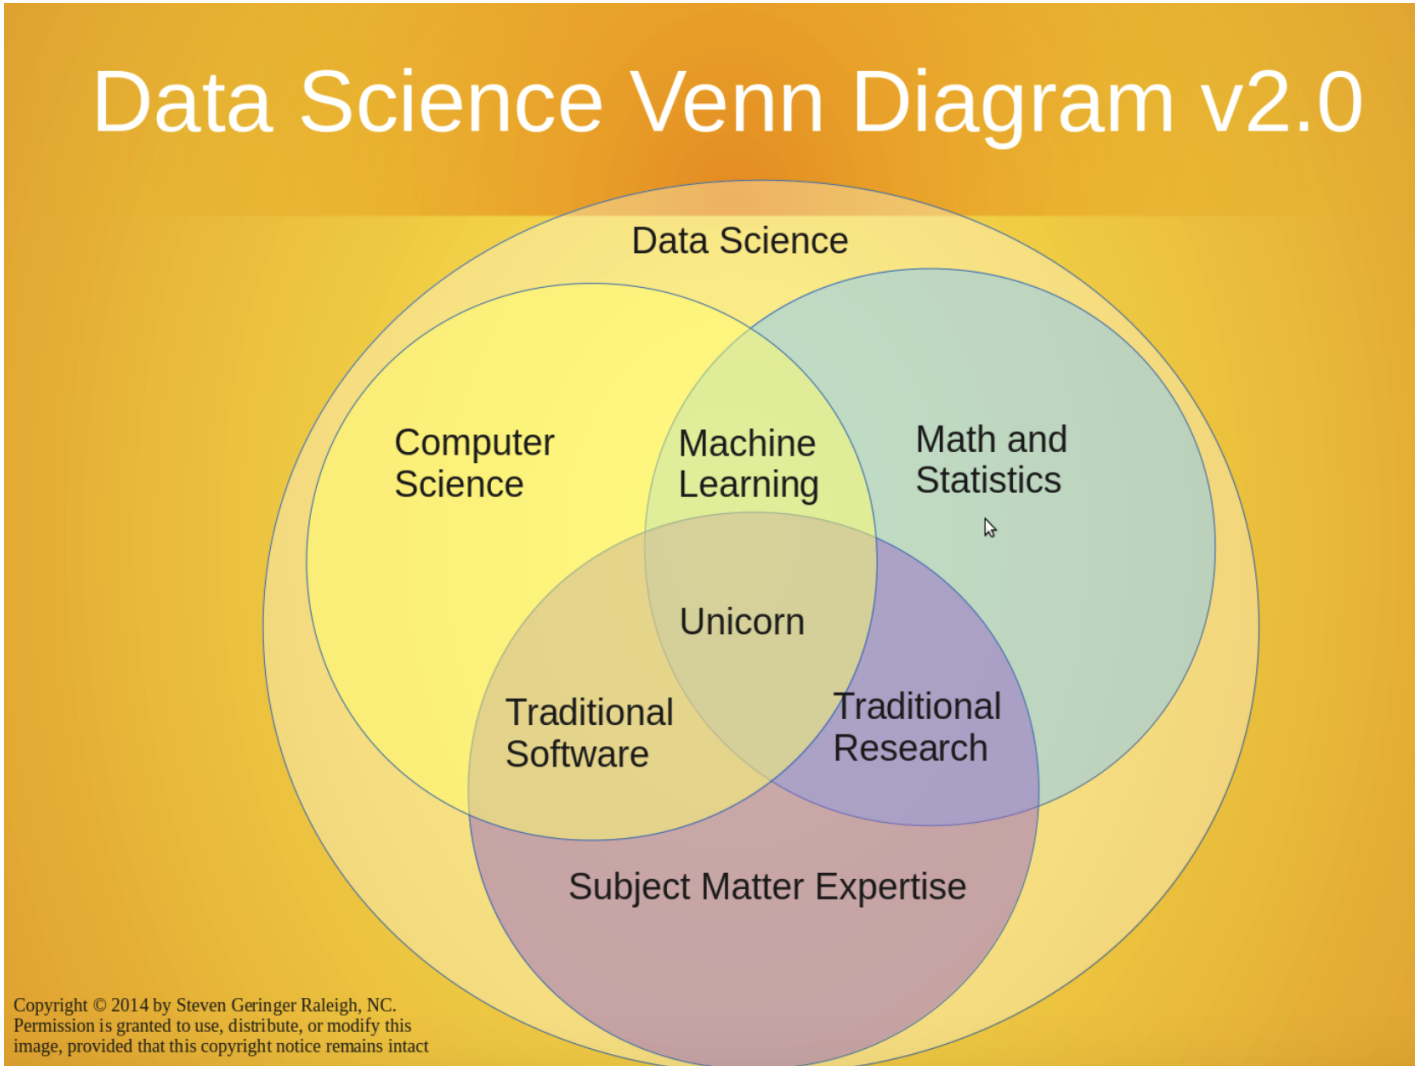

Data Scientists: The unicorn industries want?

- The field is inherently interdisciplinary
  - mathematical statistics
  - computer science
  - domain expertise
- The magical Unicorn: having all three skills
  - In 2014, these jobs go unfilled for 6 months or longer on average
- Has lead to the development of data science *teams*
  - hope is to merge skills of analysts

The Analysts of Data Science

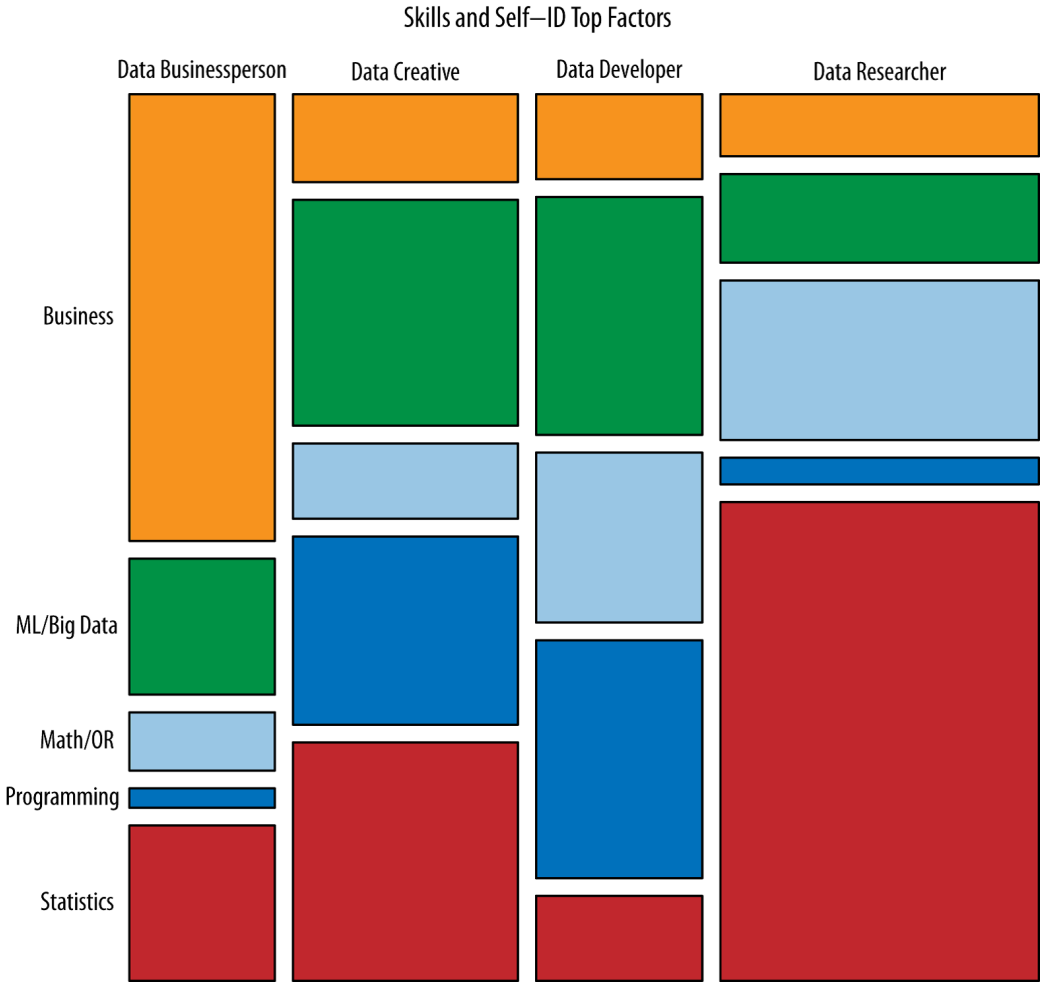

"Analyzing the Analyzers (2013) by Harry, Murphy, and Vaisman."

Software: A Data Scientist's first weapon

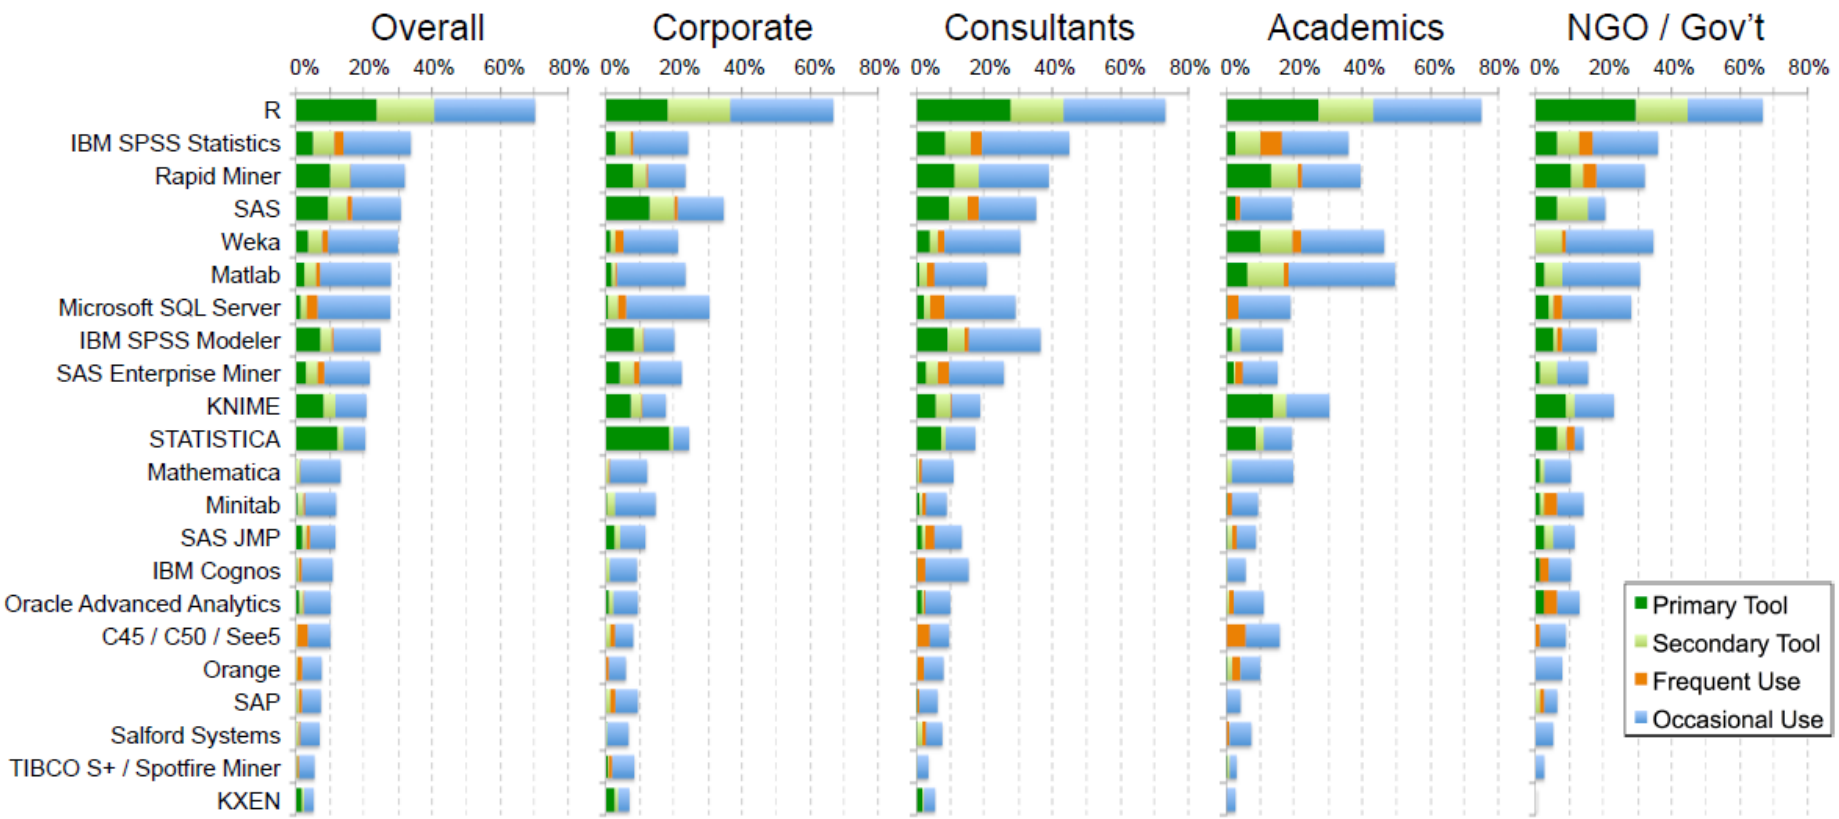

-www.datasciencecentral.com

A Data Scientist's Toolkit

Harvard's data science toolkit:

1. **Wrangle the data**: gather, clean, and sample data
2. **Manage the data**: access big data quickly and reliably
3. **Explore the data**: to make a hypothesis
4. **Make predictions**: statistical methods
5. **Communicate the results**: visualization, presentations, summaries

Get Involved! Great Resources

- Flowingdata.com
  - Contemporary visualization and data manipulation techniques
- Kaggle.com
  - Kaggle competitions: win money for solving problems!
- Coursera.org
  - Free online courses in data science and machine learning
  - 972 courses. Great resource for coding, data analysis, etc.
  - Recent notable course: "The Data Scientist's Toolbox."